

CANDATAPY

**A System for download of Cansim Data
Alpha Version 0.25**

Harold Henson

1 Introduction

Analysts working in large organizations frequently must maintain databases that support libraries of analytical tables. Challenges frequently occur in keeping these databases up to date and coping with changes in the structure of the data. Historically, teams of analysts have built substantial amounts of analytical capital with older proprietary software such as AREMOS. The collapse of budgets and the advent free data systems such as PYTHON have indicated the potential for change. CANDATAPY is a Python framework designed to allow the easy automation of the management of data provided by the Canadian government through the Cansim system.

This first version of CANDATAPY was written with a very definite application in mind. It has been documented with the full knowledge that the market of potential is limited. However, if it proves to be worthwhile, it should be possible to replicate and adapt to other situations.

The documentation starts off with a motivation of the design of the system. It is based on the authour's experience in both working on and managing teams of analysts that center their work around large systems of databases. For this reason, this documentation provides a longer section on motivation than is usual. It is hoped that the discussion of rationale for the design of this system will help a potential user better decide whether the system is suitable for their application.

After the section on motivation, there is a section that provides an overview of the use of the system. On seeing the role played by the major pieces of the system the technical details should make sense. As well, the individual components are generously documented.

2 Overview

Many analysts work in situations where they must provide complex up to date analysis based on tables that have been generated on the fly. This implies that a piece of analysis with tables and graphs must be completed in periods of time as short of a few hours. This is only possible if the analyst can safely assume that the data in the system is correct. In turn this is only possible when the software that maintains the system is stable.

CansimPY is designed to capture the data as it becomes available and incorporate it into a Pandas database. An overview of the workflow would be:

- Notification of an update to a Stats Can Matrix
- Analyst downloads the matrix from Statistic Canada web site using the sites interactive software
 - Attempts to replicate previous effort to the extent possible
- Matrix update command incorporates the new data into the central Panda database

- Programs can be run from this central database to update various BI products
- Previous version of database saved and summaries of the changes are produced

The system will work in identical fashion if it is used for a single user or for a group of analysts who use a common pool of data and need to be consistent. However, the single user can use the same system and skip some of the steps.

In software terms, this should be thought of as a framework. It consists of a few modules that provide an environment for the easy maintenance of a large database of Cansim data. Analysts should be able to maintain a Pandas database by downloading the Cansim data, and rerunning Python programs.

2.1 Motivation for Design

Many analyst are primarily subject matter experts, yet data can be the lifeblood of their work. Frequently, they work in teams that are not sufficiently large enough to justify a pure database expert. In such cases, the teams are frequently composed of a variety of individuals with varying levels of expertise in database management. In such a circumstance, the responsibility for maintaining the integrity of the collective data typically rests with one or two individuals who have the motivation and expertise.

This system is based in Python, on the assumption that although it is a programming language, it is easy enough for senior analysts to use. For organizations where this is not true, then obviously another solution is necessary. Still, it is anticipated that the use of Python by organizations who need to maintain data systems to support analysis relevant to an organization will only increase, at the expense of cost proprietary system that are not any easy to use when there is a significant amount data wrangling required to get the data in a useable format.

2.2 Target Users

The target users of the system are two-fold. First there are those who must execute the manual downloads of the Cansim data, then run the Python programs to incorporate the data into the central database. Their understanding of the system can be limited. It is only necessary for them to follow the examples in the tutorials to perform the necessary functions.

The bulk of this document targets the second class of users who assume overall responsibility for the data in their team. This class of user will first install the system on a common drive for their team. Then modify the system as requirements changes. This will include:

- Augmenting the list of Matrices from which data is downloaded
- Dropping data no longer needed from existing system

- Changing the structure of directories for which the data is kept

A serious attempt is also being made to ensure that the code is well documented, so as to complement this documentation as well.

2.3 Major Components

The system is loaded in a Python module known as CandataPY. Once this is installed on the system, all team members should be able to use the command `import CandataPY` from their account. This system will contain three major components. They will be briefly summarized in this overview section, but the section on System Components will document the system in considerable detail.

2.3.1 CanDataPY

This is the module that will be imported by the user to define their session. There are three important methods. A setup program is available to install the directory structure and files to allow the system to run. It should only be run once.

2.3.2 CanDataMatrix

A substantial portion of setting up the system will occur at this stage. A separate class for each Stats Can Matrix will be prepared. These classes will inherit from the generic StatCanMatrix class. The manipulations specialized to this data will be included in the inherited class. For example if industry specific data needs to be aggregated in a certain fashion, then expressed as shares, it could be done at this stage.

2.3.3 TimeVar.py

This object contains the finest level of detail. When a matrix is updated each one of these objects will be invoked from TSlist. It will read the data, and then compare it with what is already in the Pandas database. If the data is not within an acceptable level of tolerance, the version on the Pandas database is not changed and a warning is issued.

3 Use

This module should be thought of as a framework built in Python. Once the startup program has been run, the analysts will create a program for each matrix that can be run to add the data from the matrix to a central database. These programs can be saved and run on an as needed basis.

3.1 Run startup program

The first step is to establish a directory in the network shared by all members of the team. Analysts who are maintaining the data base should have full access to this directory and its sub-directories. All possible users should have at least read access to the directory where the final Pandas database should reside.

The second system is to install the CansimPY module on the "System", so that all analysts maintaining the database will be able to import the module when they run a Python program. As this is very site specific, not much will be said in this document. However, it should be no different than installing any other shared module for team use.

After that go into Python, and set the working directory to the target directory. Then from the Python console enter the command `CansimPY.setupdirs()`. This will setup all the directories. In general it is advised not to change any of the defaults.

For the subsequent discussion the root directory is referred to as Cendir. Sub-directory will be designated with periods to avoid confusion with the direction of the slash.

3.2 Download the data

It is recommended that the analyst download the data from Statistics Canada down to a local drive. Then they can go into their spreadsheet viewer and look at the data to quickly verify by means of comparison of the previous download of that matrix, which resides in Cendir.archives. The bulk of the possible problems can be identified with this visual inspection. As confidence in the system grows, the analyst may skip this step and rely on the automated procedures.

If the analyst is confident that he has downloaded the data correctly then the freshly downloaded data should be copied down into Cendir.rawdownloads for further processing. The analysts should not be overly concerned that they may have overwritten good data as Cendir.archives contains the downloaded data that is in use with the current Pandas database.

There are many different choices that must be made when downloading the data. Currently, the analyst is restricted in the number of styles of download the software can accept. As this product matures, it will be able to manage a wide variety of downloads, which in turn will make things easier for the average user.

This section outlines in thematic terms the choices confronted by the analysts. In each case the implications of these choices are discussed. The implications of these choices are discussed in the next section on loading the matrix. It is appreciated that many analysts will become impatient with just a theoretical discussion. Thus a tutorial will be provide a step by step process.

The choices can be summarized along these thematic lines:

3.2.1 Data to be downloaded

The matrix is the logical organizing unit for the data retrieval. Unless there are complicating factors, most analysts will find it optimal to download all the data, which is done by checking the "all" boxes in the select tab. Once the data is downloaded, then the analyst would run statistical programs that would call in the data that is needed to produce their reports and tables.

As some matrices contain thousands of time series, it is likely that much of the data is never used. However, given the cost of storage are so miniscule there is likely no harm in downloading too much data. However, it is appreciated that some users may find themselves in organizational environments where this would be thought to be wasteful. If this is the case, options are available to reduce the volume of data.

For a start, the Cansim downloading program can offer very useful options for downloading subsets of the data. For example, if only data at the national level is required, it is often possible to suppress provincial detail and reduce the size of the downloaded data by a factor of almost ten.

It is appreciated that it may be difficult to achieve consistent selection of these options on a regular basis. As a result, some options are available to avoid the downloading of the entire matrix. These are:

- List of vector numbers - use method `set_list` to pass a list of strings of v numbers that will only be included.
- Range of vector numbers to be included - Use the `set_incrng` to specify the lowest and highest vnumber in the desired matrix. Note the lowest should come first.
- Range of vector numbers to be excluded - Use `set_exrng` in a fashion similar to `set_incrng`. Note that in both these cases only two v numbers should be specified.

3.2.2 Organization of data

Each matrix can contain a vast amount of the data, even within the structure known as CSV. There are a number of key choices:

- time as column or rows - It is recommended that time as rows be chosen.
- inclusion of vnumbers - It is essential that the vector number be included for the system to work.

For this system to work, it is important to always include the vnumbers when prompted towards the end of process. The numbers are unique to each time series and are used to identify the data when it is exported to other systems.

The data when it is downloaded into the csv format will be viewed as a two dimensional matrix. Basically, the structure is for the most part defined by whether the time dimension is along the horizontal or vertical axis. Two

examples are displayed in the Annexes 7.4.1 and 7.4.2. The system defaults to time along the vertical axis but in repeated patterns for each variable, as is shown in Annex 7.4.1. The reasons for doing this are technical and historical, as it has been awkward to work with records that are extremely wide, which is what would happen with time on the vertical axis and thousands of time series on the horizontal.

If the user downloads the data with time on the horizontal axis, it is necessary to specify `'by_row'` in the upload command rather than `'obs_by_row'` as specified in the example. It is also important to note that it is not necessary to date column as it will be assumed to be the same as the first column of data.

It should be noted that at this early stage that not all possible configurations of data are even mentioned. However, the two formats specified in the annex have been tested and known to work.

3.2.3 format of the data

This package only supports data downloaded into the CSV format. The package will not support the other two choices, Beyond 20/20 or XML. Note that it is presumed that the data will assumed to not have been manipulated and the "Data as R

3.3 Load Matrix of data into Central Database

There will be a CandataPy object underlying every matrix. The analysts will simply execute the update method for that matrix. This should fit into the workflow of the group in a natural way as all the data from any given matrix will typically come from the same source and be updated at the same time.

Most of the time, things will work with very little effort. However, the system may produce warnings that need to be addressed. A file `session_log.txt` will be produced that will list any variables that could not be produced.

At this point in time a decision may have to be made about whether the data needs to be rethought. Many times the warnings can be ignored but other times adjustments must be made. It is important to note that these decisions can only be made by a subject matter expert, as it is not a technical issue.

3.3.1 Define Frequency of Data

Every matrix will share a common frequency. However, the start date and sometimes the end date may vary. As a result, the package will adjust the start date anew for each time-series. However, the frequency is defined when the object is created for the matrix. This is done by passing the frequency into the inherited version of upload from the parent StatCanMatrix class.

In the example in the Annex 'M' is used as the second argument to defined the monthly data. This comes from the list of designators for frequencies of time given on page 296 of McKinney. This list is far more detailed than

necessary for Statistics Canada Cansim data. At this point in time the following frequencies have been implemented:

- - 'M' for monthly data
- - 'AS-JAN' for annual data starting at the beginning of January.
- - 'Q-MAR' for quarterly data with first quarter ending in March.

Some analysts may be surprised with the richness of definitions offered in McK-inney. This is a result of Pandas rethinking a lot of these concepts from first principles. In the case of quarterly data, it is described a 'Q-MAR' rather than just Q1. This designation comes from the Panda's description and refers to quarterly data where the first quarter ends in the month indicated. This lines up with what is seen on the CSV file as the first quarter of a year is designated with the digit 3 after the slash.

An important practical implication of this is that if an analysts wants to refer to a single observation of quarterly data they would want to use the last month in the quarter rather than the number of the quarter itself. This point is being stressed as it is a likely source of errors. However, in Pandas an index of 2015-3 refers to the first quarter of 2015 and not the third quarter.

3.3.2 Define File Name

At this point, the analyst should have a version of the downloaded file in the subdirectory rawdump. The exact name of the file needs to be entered when creating the object specific to the Stats Can Matrix. This is shown in Listing 1 where this282 is created.

3.4 Keep data up to date

Once a matrix of data has been added to the central database, the source code will be saved in the archive directory. At a future date, this code can be rerun without significant modification so as to update the database.

3.5 export data

It is certainly possible to build reports off this system. In fact, a demo of this system will be built there the database drives some Django web tables. However, it is recognized that much analysis will occur with the more specialized data that exists. As such, all of the export capacity that exists for a Pandas database still exists.

3.6 Quick Start Summary

The above explanation can be summarized as:

- Select empty directory on Shared Drive

- Copy three Python Files into Directory
- Go into Python Version 3 (Anaconda - Spyder preferred)
- Run setup script in CansimPY.PY
- Validate that subdirectories are created as described in Annex
- Download Statscan Matrix by rows with headers in CSV format
 - At present only monthly data will work
- Place CSV file in rawdump director
- Modify example in Annex
- run script with upload method
 - Do not modify the three existing files but import them into your script
 - May have to pip some libraries
- Validate data in h5 file

4 System Components

This section is written as a reference and will be frequently updated. Its goal is support the day to day use, rather than explain the system. This section may form a separate document if it becomes disproportionately large relative to the rest of the document.

4.1 environment

It is assumed that all possible user's of the data will have access to a common directory. Through out this document, this common directory will be referred to as cendir, although the actual name will be determined by the local environment. A period will be used in place of a slash when discussing sub-directories, as the direction of the slash may vary with the operating system.

In this directory a file central_data.h5. This file type supports a write once read many times usage. It is thus expected that all potential data users will have access to this file. Those few who will maintain the database should have read and write access. When the installation is being executed by the startup program, it should be possible to create subdirectories in cendir.

The system is written to favour the recent versions of Python 3. As Pandas is sometimes difficult to install on some environments, Anaconda is often the simplest solution as Pandas is part of the suite. There may be prompts for other packages, but they can be resolved with PIP.

4.2 modules

The three basic modules were introduced in the Overview section.

4.2.1 CanDataPYSession.py

This module has two roles. First, it sets up the directory structure in Cendir. There is a startup.py in this module that is run once. Second it provides a parent class that is inherited by StatsCanMatrix. This parent class maintains information common to all updates of the matrices.

- Central.data - This is the file handle to .data.h5 in the main directory.
- cwd - The drive where the system is installed.

4.2.2 StatsCanMatrix.py

The data from Statistics Canada web site will be downloaded in groups referred to as as matrices. In general, individuals will download a matrix that may or may not contain every possible Timeseries from the matrix. Still, the Cansim matrix constitutes a fundamental organizing principle around the data. As such, this system has a corresponding object that stores the data. The information in this matrix is persistent, which allows it to be updated on a regular basis. The StatCanMatrix object will be the parent of the objects that are created for each Stats Can matrix.

The definition of each new class for each download to occur will contain information particular to that matrix. An example is provided in the sample listing at the end of this document. In this listing the user provides four pieces of key information. These are sufficient to allow the loading of the downloaded CSV file into the Panda file.

- Update - This method will attempt to load the downloaded csv file. It will return True if the data is successfully downloaded and converted to the Pandas format. The results of this process will be stored in a dictionary update_info.
- __str__ This is overridden to allow the quick view of the update_info dictionary
- filename - This is string that contains the name of the Statscan that has been downloaded and is located in the cendir.rawdump directory.
- ses_log - A file handle to the log file that is maintained for just the session. This will be very detailed in nature.
- SCfilehandle - This is the handle to filename.

4.2.3 timevar.py

The Panda's Series is the fundamental building block of the system. The module timevar.py defines a class, which builds up the data to create a Panda Series. When all of the data has been read in, this data is saved to a dataframe that it has been saved to. This process is repeated until all the desired data for the matrix has been added to thematrix.thepandas.

This class exploits the fact that some aspects of the series are the same for all the series on a Matrix. As such, there is one initialization program for all of the parameters, and there is a second for just the variables that may change. Detailed documentation is embedded in the source code itself, however and overview of the key methods follows:

- `__init__` - The frequency is established here. As well, a link is formed to the larger session with the thematrix. Note that a method is used to initialize the various values, which in turn calls two different routines.
- `setdate(thetoken)` - This method will take a token and parse out it according to the frequency that was established in `__init__`. The `datstr` is created so that it can be used by `save` to create the series
- `setvalue(thetoken)` - This method converts thetoken into a float and appends it to the values list. If the thetoken cannot be converted to a real number, it will be converted to `numpy.nan` if it is '..', '...', or 'x'. Any other value it is assumed that there is some kind of processing error and the entire series is not used. The name of the unused series is written out to the error log. See Appendix for some details on this issue.
- `save` - When the last observation has been read in, this method will be called to add the time series to the dataframe. For this method to work properly:
 - the beginning period needs to have been established
 - the number of observations should be known
 - there should be a vector the length of the number of observations
 - the name of the variable should have been established

5 Future Directions

This version contains a very minimal capacity. While a reasonable as a first step, there is still much than be done to fully automate the process of automating the production of Business Intelligence with Cansim data. As it stands, there are three major areas of potential development:

- **Resolve Obvious Shortcomings** - The current Alpha version can only accept monthly data downloaded by row. This will be relaxed in the immediate future.

- Automate Download - Currently the assumption is that analysts will use the Stats Can interactive environment to download in the Matrix. Hopefully, the updates will follow the same pattern every time the database is updated. However, this process is inefficient and error-prone in comparison to a possible automated process.

It is technically possible to automate this process. However, it is important to note that Statistics Canada's position on this has been in flux in recent history. Still, a utility to replicate the download process may be very useful at some point in the future.

- Quality Control of Data - It is impossible to manually verify all the after each download. However, if comparisons could be done with a version that had been manually verified, then many possible problems could be captured. A section of this is included in Version 0.1 of this data, but has been omitted for now.
- Automated tables in Django - It would also be very possible to automated the production of reports in Django for Intranet/Internet sites. Not only would the tables be useful in themselves, they would also provide an opportunity to blend narrative data that helps understand the quantitative data coming from Statistics Canada.

6 References

McKinney Wes, (2013), "Python for Data Analysis", O'Reilly.

7 Annex

7.1 Treatment of Invalid Data

Data downloaded from Cansim should be numeric. However, if there are problems Statistics Canada will supply text in place of the numbers. To date, three text items have been identified that flag reasons that an numeric value of an observation is not available:

- 'x' - Data suppressed to meet confidentiality requirements.
- '..' - Data not available for the reference period
- '...' - figures on applicable In all three of these cases, the Pandas value for NA is inserted in the observation. Note that this usage of is defined in page 140 of McKinney.

7.2 Code Samples

The box gives a working example that loads Matrix282001 into Pandas database.

Listing 1: A matrix inherited class example

```

class Matrix282_0001(StatCanMatrix):
    def __init__(self, thefile, thename="282"):
        super(Matrix282_0001, self).__init__(thefile, 7000,
                                              matname=thename)
        # these are valid keys
        self.setcollist(['date', 'NA', 'NA', 'gender',
                        'NA', 'VName', 'NA', 'datum'])
    def upload(self):
        super(Matrix282_0001, self).upload('obs_by_row',
                                             'M', 'first5')
# create an instance of it
this282 = Matrix282_0001("02820001-eng.csv", 2820001)
# to the actual upload
this282.upload()

```

7.3 File Structure

As stated in the text, the CanData files can be setup anywhere on the shared drive where various members of the team can get access. To start a session they need to be able to set this drive, referred to as Cendir above, as the current working directory. To verify the setup.py has installed the system properly, the following files should be present

Note that periods rather than slashes will be used to indicate the subdirectories. This is because with Dos derived operating systems the slash is opposite to the Unix variants.

7.3.1 cendir

This is the working directory. In practice its name may be quite long. The analysts who are responsible for updating the data should have full access to it. There should be only two non-python files in this directory besides the subdirectories listed below. However, it is realized that often the base directory becomes cluttered. With this documentation, a supervisor of this process can clean out extraneous files with confidence.

The two files are Central_data.h5 and Central_log.txt, which are created by the installation program. These two files should be backed up on a regular basis, although as will be seen there will be copies in the archive directory. However, it is possible that the system source code may left in this directory if there are problems with the import function. If this is the case, there may be the three python files that make up the system in this directory.

7.3.2 cendir.rawdump

The system will look in this directory for the raw data. Once the file has been processed, it will be deleted. It is assumed that there will be no cases where more than one analysts downloads files with exactly the same name.

7.3.3 cendir.archive

Every time an analyst updates Central_data.h5 with new data copies of the old log file are archived in a unique subdirectory that is created in cendir.archive. As well, there will be a detailed run specific log file kept as well a copies of matrix-specific version of the pandas database. This allows for complete diagnostics if problems occurs.

This may seem like overkill to some. It is likely that some sort deletion of old directories will occur. This is entirely optional as once the Central_data.h5 file has been updated, these archives have no functional value.

The use of this directory will largely be dictated by the culture of the team using the system. Often, very low tolerance for errors characterizes the work of many of the users of this kind of data. As a result, the desire to perform forensic reviews to diagnose the error source can be fairly high. The option of providing user names is provide a potential to conduct very detailed historical reviews.

7.4 Sample File Structures

An important component of making the system is work is the downloading of the files into a consistent structure so that it can be important. Much of the discussion in the text will seem abstract without experience with the different possible structures. This section provides screenshots of the file various file structures with the goal of providing support to the discussions in the main text.

7.4.1 oneobperrow

This was the first type supported by the software, and is the default.

A	B	C	D	E	F
Ref_Date	GEO	SECTOR	Vector	Coordinate	Value
2009	Canada	Total, industries and households	v79874995	1.1	721165
2010	Canada	Total, industries and households	v79874995	1.1	735016
2011	Canada	Total, industries and households	v79874995	1.1	742181
2012	Canada	Total, industries and households	v79874995	1.1	747471
2013	Canada	Total, industries and households	v79874995	1.1	764223
2014	Canada	Total, industries and households	v79874995	1.1	768238
2009	Canada	Total, industries	v79874996	1.2	584244
2010	Canada	Total, industries	v79874996	1.2	596356
2011	Canada	Total, industries	v79874996	1.2	599039

Note that each year is a different value for the same Vector number.

7.4.2 byrow

There is one v number per row.

1	Table 153-0114	Physical flow account for greenhouse gas emissions, annual (kilotonnes)(1,2,3,4,5,9,10,11,14)			
2	Survey or program details:				
3	Canadian System of Environmental-Economic Accounts - Physical Flow Accounts - 5115				
4	Geography	Sector	Vector	Coordinate	2010
5	Canada	Total, industries and households	v79874995	1.1	735016
6	Canada	Total, industries	v79874996	1.2	596356
7	Canada	Crop and animal production [BS11A00]	v79874997	1.3	69899
8	Canada	Forestry and logging [BS11300]	v79874998	1.4	8758
9	Canada	Fishing, hunting and trapping [BS11400]	v79874999	1.5	590